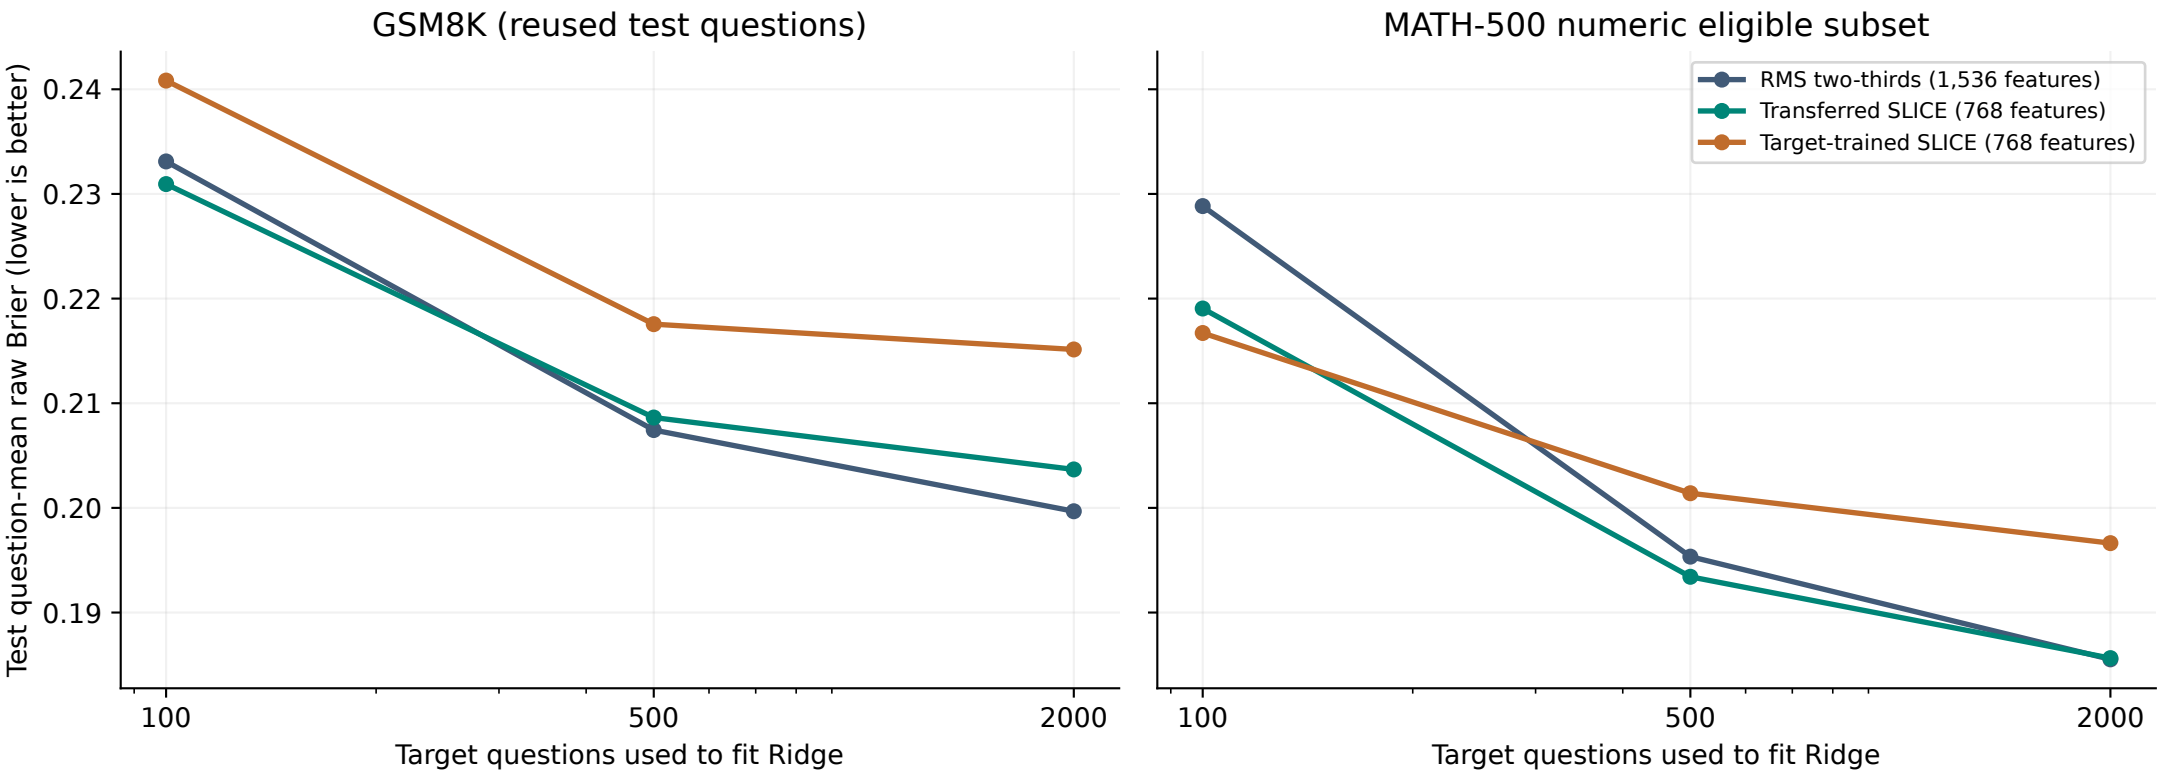
Feature transfer: prediction versus target head-fitting data
GSM8K (reused test questions)
MATH-500 numeric eligible subset
Test question-mean raw Brier (lower is better)
0.24
0.23
0.22
0.21
0.20
0.19
100
500
2000
Target questions used to fit Ridge
RMS two-thirds (1,536 features)
Transferred SLICE (768 features)
Target-trained SLICE (768 features)
Offline prediction; frozen Qwen. Five compressor seeds share generated answers. Inspected data; exploratory results.
Points are seed means; differences and uncertainty are shown separately.